

AAA
BBB
CCC
DD

Lines: 4

Acknowledgments

We are grateful to the many colleagues whose insights have shaped this work, often through brief but illuminating conversations that deepened our understanding of machine learning and artificial intelligence. These seemingly small exchanges in hallways, at conferences, and over coffee have had profound impacts on our thinking about instruction-following in language models. We also thank the model providers and the OpenRouter platform for enabling this comprehensive evaluation across diverse architectures.

References

1. Y. Chen, X. Lu, Y. Zhou, Y. Liu, X. Wang, and D. Chen. Sifo: A benchmark for sequential instruction following in open-domain settings. *arXiv preprint arXiv:2405.09723*, 2024.
2. Q. He, L. Liao, Z. Chen, Y. Zhang, H. Dong, H. Li, Y. Liu, X. Chen, and D. Li. Cello: Chinese complex instruction following benchmark. *arXiv preprint arXiv:2410.13382*, 2024.
3. D. Hendrycks, C. Burns, S. Basart, A. Zou, M. Mazeika, D. Song, and J. Steinhardt. Measuring massive multitask language understanding. *arXiv preprint arXiv:2009.03300*, 2021.
4. Y. Jiang, Y. Wang, X. Zeng, W. Zhong, L. Li, F. Mi, L. Shang, X. Jiang, Q. Liu, and W. Wang. Followbench: A multi-level fine-grained constraints following benchmark for large language models. *Annual Meeting of the Association for Computational Linguistics*, 2023.
5. S. Kim, J. Suk, S. Longpre, B. Y. Lin, J. Shin, S. Welleck, G. Neubig, M. Lee, K. Lee, and M. Seo. Prometheus 2: An open source language model specialized in evaluating other language models. In *Conference on Empirical Methods in Natural Language Processing*, 2024.
6. N. Lambert, V. Pyatkin, J. Morrison, L. Miranda, B. Y. Lin, K. Chandu, N. Dziri, S. Kumar, T. Zick, Y. Choi, et al. Rewardbench: Evaluating reward models for language modeling. *Neural Information Processing Systems*, 2024.
7. J. Li, S. Sun, W. Yuan, R.-Z. Fan, H. Zhao, and P. Liu. Generative judge for evaluating alignment. *International Conference on Learning Representations*, 2023.
8. X. Li, P. Yu, C. Zhou, T. Schick, L. Zettlemoyer, O. Levy, J. Weston, and M. Lewis. Self-alignment with instruction backtranslation. *International Conference on Learning Representations*, 2023.

-
9. Y. Liu, K. Shi, A. R. Fabbri, Y. Zhao, P. Wang, C.-S. Wu, S. Joty, and A. Cohan. Reife: Re-evaluating instruction-following evaluation. *North American Chapter of the Association for Computational Linguistics*, 2024.
 10. S. Oh, J. Thorne, and A. Vlachos. Instructir: Content-level instruction following for information retrieval. *arXiv preprint arXiv:2408.05360*, 2024.
 11. F. M. Polo, L. Cisneros-Velarde, N. Kannan, Y. Cheng, H. He, K. C.-C. Chang, J. Sun, and Y. Zhang. tinybenchmarks: evaluating llms with fewer examples. *arXiv preprint arXiv:2402.14992*, 2024.
 12. Y. Qin, K. Song, Y. Hu, W. Yao, S. Cho, X. Wang, X. Wu, F. Liu, P. Liu, and D. Yu. Infobench: Evaluating instruction following ability in large language models. *Annual Meeting of the Association for Computational Linguistics*, 2024.
 13. M. T. Ribeiro, T. Wu, C. Guestrin, and S. Singh. Beyond accuracy: Behavioral testing of nlp models with checklist. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 4902–4912, 2020.
 14. O. Skopek, R. Aralikatte, S. Gooding, and V. Carbune. Towards better evaluation of instruction-following: A case-study in summarization. In *Conference on Computational Natural Language Learning*, 2023.
 15. A. Srivastava, A. Rastogi, A. Rao, A. A. M. Shoen, A. Abid, A. Fisch, A. R. Brown, A. Santoro, A. Gupta, A. Garriga-Alonso, et al. Beyond the imitation game: Quantifying and extrapolating the capabilities of language models. *arXiv preprint arXiv:2206.04615*, 2022.
 16. Y. Wang, H. Ivison, P. Dasigi, J. Hessel, T. Khot, K. R. Chandu, D. Wadden, K. MacMillan, N. A. Smith, I. Beltagy, et al. How far can camels go? exploring the state of instruction tuning on open resources. *Neural Information Processing Systems*, 2023.
 17. Y. Wang, Y. Kordi, S. Mishra, A. Liu, N. A. Smith, D. Khashabi, and H. Hajishirzi. Self-instruct: Aligning language models with self-generated instructions. *arXiv preprint arXiv:2212.10560*, 2022.
 18. Y. Wang, Z. Yu, Z. Zeng, L. Yang, C. Wang, H. Chen, C. Jiang, R. Xie, J. Wang, X. Xie, et al. Pandalm: An automatic evaluation benchmark for llm instruction tuning optimization. In *International Conference on Learning Representations*, 2023.
 19. J. Wei, X. Wang, D. Schuurmans, M. Bosma, E. Chi, Q. Le, and D. Zhou. Chain of thought prompting elicits reasoning in large language models. *arXiv preprint arXiv:2201.11903*, 2022.
 20. O. Weller, M. Marone, D. Lawrie, I. Mackie, L. Soldaini, and B. Van Durme. Followir: Evaluating and teaching information retrieval models to follow instructions. *Annual Meeting of the Association for Computational Linguistics*, 2024.
 21. B. Wen, P. Ke, X. Gu, L. Wu, H. Huang, J. Zhou, W. Li, B. Hu, W. Gao, J. Xu, et al. Benchmarking complex instruction-following with multiple constraints composition. *Neural Information Processing Systems*, 2024.
 22. C. Xu, Q. Sun, K. Zheng, X. Geng, P. Zhao, J. Feng, C. Tao, and D. Jiang. Wizardlm: Empowering large language models to follow complex instructions. *International Conference on Learning Representations*, 2023.

-
23. Z. Xu, F. Jiang, L. Niu, Y. Deng, R. Poovendran, Y. Choi, and B. Y. Lin. Magpie: Alignment data synthesis from scratch by prompting aligned llms with nothing. *arXiv preprint arXiv:2406.08464*, 2024.
 24. Z. Zeng, J. Yu, T. Gao, Y. Meng, T. Goyal, and D. Chen. Llmbar: Evaluating complex natural language understanding with llms. *Conference on Empirical Methods in Natural Language Processing*, 2024.
 25. J. Zhou, T. Lu, S. Mishra, S. Brahma, S. Basu, Y. Luan, D. Zhou, and L. Hou. Instruction-following evaluation for large language models. *arXiv preprint arXiv:2311.07911*, 2023.
 26. K. Zhou, Y. Zhao, Y. Liu, K. Song, W. X. Zhao, and J.-R. Wen. Instruction-aware dense retrieval with self-knowledge distillation. *Information Processing and Management*, 2024.
 27. Y. Zou, X. Chen, K. Lu, L. Zhang, and Q. Liu. Eifbench: A benchmark for extremely complex instruction following. *arXiv preprint arXiv:2501.08901*, 2025.